

```

In [1]: import pandas as pd
        from sklearn.preprocessing import StandardScaler
        from sklearn.model_selection import train_test_split
        from sklearn.linear_model import LogisticRegression
        from sklearn.ensemble import RandomForestClassifier
        from sklearn.neural_network import MLPClassifier
        from sklearn.metrics import confusion_matrix, classification_report, roc_auc

/Users/manikshakya/anaconda3/envs/hgp/lib/python3.10/site-packages/scipy/_i
nit__.py:146: UserWarning: A NumPy version >=1.17.3 and <1.25.0 is required
for this version of SciPy (detected version 1.25.0
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")

In [2]: # Read the dataset
df_train = pd.read_csv('mimic_train.csv')
df_test = pd.read_csv('mimic_test.csv')

In [3]: # Check for Null values and remove the unnecessary rows from training and te
df_train.drop(['SUBJECT_ID', 'HADM_ID', 'ADMITTIME', 'DISCHTIME', 'DAYS_NEXT
              'ADMISSION_TYPE', 'DEATHTIME', 'DISCHARGE_LOCATION', 'INSURAN
              'DIAGNOSIS', 'GENDER', 'DOB'], axis=1, inplace=True)
df_test.drop(['SUBJECT_ID', 'HADM_ID', 'ADMITTIME', 'DISCHTIME', 'DAYS_NEXT_
              'ADMISSION_TYPE', 'DEATHTIME', 'DISCHARGE_LOCATION', 'INSURANC
              'DIAGNOSIS', 'GENDER', 'DOB'], axis=1, inplace=True)

In [4]: # Transform non-numerical features into categorical type
categorical_cols = ['blood', 'circulatory', 'congenital', 'digestive', 'endo
                  'infectious', 'injury', 'mental', 'misc', 'muscular', 'n
                  'pregnancy', 'prenatal', 'respiratory', 'skin']

df_train[categorical_cols] = df_train[categorical_cols].astype('category')
df_test[categorical_cols] = df_test[categorical_cols].astype('category')

In [5]: # Select only important features for training and testing purpose
selected_features = ['blood', 'circulatory', 'congenital', 'digestive', 'endo
                  'infectious', 'injury', 'mental', 'misc', 'muscular', '
                  'pregnancy', 'prenatal', 'respiratory', 'skin']

X_train = df_train[selected_features]
y_train = df_train['OUTPUT_LABEL']

X_test = df_test[selected_features]
y_test = df_test['OUTPUT_LABEL']

In [6]: # Re-scale the data using StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

In [7]: # Train-test split
X_train_scaled, X_val_scaled, y_train, y_val = train_test_split(X_train_sca

In [8]: # Train the models
models = [
    LogisticRegression(),
    RandomForestClassifier(),
    MLPClassifier()
]

```

```
In [9]: for model in models:
        model.fit(X_train_scaled, y_train)
        y_pred = model.predict(X_val_scaled)
        y_pred_prob = model.predict_proba(X_val_scaled)[: , 1]

        # Evaluate the model
        print(f"Model: {type(model).__name__}")
        print("Confusion Matrix:")
        print(confusion_matrix(y_val, y_pred))
        print("Classification Report:")
        print(classification_report(y_val, y_pred))
        print("AUC Score:", roc_auc_score(y_val, y_pred_prob))
        print("-----")
```

Model: LogisticRegression

Confusion Matrix:

```
[[142  46]
 [ 95 117]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.60	0.76	0.67	188
1	0.72	0.55	0.62	212
accuracy			0.65	400
macro avg	0.66	0.65	0.65	400
weighted avg	0.66	0.65	0.64	400

AUC Score: 0.7249849458048976

Model: RandomForestClassifier

Confusion Matrix:

```
[[121  67]
 [ 65 147]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.65	0.64	0.65	188
1	0.69	0.69	0.69	212
accuracy			0.67	400
macro avg	0.67	0.67	0.67	400
weighted avg	0.67	0.67	0.67	400

AUC Score: 0.7333651144118828

Model: MLPClassifier

Confusion Matrix:

```
[[120  68]
 [ 85 127]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.59	0.64	0.61	188
1	0.65	0.60	0.62	212
accuracy			0.62	400
macro avg	0.62	0.62	0.62	400
weighted avg	0.62	0.62	0.62	400

AUC Score: 0.6680048173424327

```
/Users/manikshakya/anaconda3/envs/hgp/lib/python3.10/site-packages/sklearn/neural_network/_multilayer_perceptron.py:692: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.  
  warnings.warn(
```

Compare the results and describe which classifier is performing better and why?

Based on the evaluation metrics, RandomForestClassifier outperforms the other classifiers in terms of accuracy, precision, recall, and F1-score. It achieves the highest accuracy of 0.67, indicating that it predicts the output label correctly for approximately 67% of the instances.

In []: